



(3 points) What is highest speedup possible through pipelining for a 6 instruction program if latch delay is 2 ns and total combinational logic delay of a non-pipelined design is 10 ns?

The highest speedup can be obtained when dividing the total combinational logic into N stages, which each stage taking $10/N$ ns.

On the non-pipelined processor, the program takes $6 \times 10 = 60$ ns to run.

On a pipelined processor, each stage takes $\frac{10}{N} + 2$ ns. The program in total takes

$$(N-1+6) \times \left(\frac{10}{N} + 2 \right) = (N+5) \left(\frac{10}{N} + 2 \right) \text{ ns}$$

The total time is minimized for $N = 5$. The speedup is $\frac{60}{30} = 2$.

(12 points) Fill in the blanks and calculate the cache access times for the following actions in sequence (the second action follows immediately after the first one). Show your calculations for full credit. Write the address in hex and circle hit or miss.

i. Read virtual address 4x00107

• Cache access time: 207 ns

• Physical address: x3307

• Cache hit / miss

• TLB hit / miss

VA: 00 0000 0001 0000 0111 TLB miss on VPN = 00 0000 0001

PA: 0011 0011 0000 0111 Cache Compulsory Miss

5 + 80 + 5 = 90 ns to get the physical tag

4 ns to detect a cache miss by tag comparison

100 ns to access the main memory

10 + 3 = 13 ns to access data (tag access can be done in parallel)

ii. Read virtual address 4x34500

• Cache access time: 13 ns

• Physical address: x3300

• Cache hit / miss

• TLB hit / miss

VA: 11 0100 0101 0000 0000 TLB hit on VPN = 11 0100 0101

PA: 0011 0011 0000 0000 Cache hit

Parallelism:

Indexing tag + tag comparison = 8 + 4 = 12 ns

Indexing data + multiplexing data = 10 + 3 = 13 ns

• Byte addressable

• 256 byte (2^8 byte) page size

Cache:

- Virtually-indexed and physically-tagged
- 4-way set-associative with 64 index bits
- 4KB (2^{12} byte) 1 KB (2^{10} byte) data storage (excluding bits for dirty, valid, tag and LRU)
- Read allocate policy
- Indexing the data array takes 10 ns
- Indexing the tag array takes 8 ns
- Tag comparison takes 4 ns
- Multiplexing the output data takes 3 ns
- A cache miss takes 100 ns to access the main memory and allocate to the cache line
- Assume a hit or miss is detected immediately after the tag comparison
- Initially empty (all lines are invalid)

TLB:

- Fully-associative
- A TLB access takes 5 ns
- Read allocate policy
- TLB is updated on a TLB miss
- All entries are listed below

Valid	VPN	PPN
0	00 0000 0001	0001 0000
1	00 0000 0010	0000 1110
0	00 0001 0110	0011 0011
0	00 0001 1011	0000 0000
1	11 1010 0100	1000 0100
1	11 0100 0101	0011 0011
1	10 0010 1010	1100 0110
0	00 0000 0000	0000 0001

Page Table:

- Single level page table
- A page table access takes 80 ns
- Some of the entries are listed below

Valid	VPN	PPN
0	00 0000 0000	0001 0000
1	00 0000 0001	0011 0011
1	00 1101 0010	0000 0000
0	01 0000 0011	0010 0100
1	01 1010 0010	1110 0001
1	01 1111 1101	0000 1110
1	10 0000 0110	1100 0110
0	11 1111 1111	0101 0110

Fast cache access

- Only require address translation upon miss

Issues

- Homonym
 - Same VA maps to different PAs upon context switch
- Synonym (also a problem in VIPT)
 - Different VAs map to the same PA when data is shared by multiple processes

```

1 add $t0, $s0, $s1
2 xor $t1, $t0, $s2
3 lw $s0, -12($s0)
4 sub $s5, $s0, $s1

```

possible to resolve any of the forwarding would be unnecessary to resolve one, but it is sub.

```

1 add $t0, $s0, $s1
2 lw $s0, -12($s0)
3 sub $s5, $s0, $s1
4 xor $t1, $t0, $s2

```

